

A dataset of software collapse in scientific science

Samuel Grayson

Univerity of Illinois Urbana Champaign
Urbana, IL
0000-0001-5411-356X

Daniel S. Katz

Univerity of Illinois Urbana Champaign
Urbana, IL
0000-0001-5934-7525

Reed Milewicz

Sandia National Laboratories
Albuquerque, NM
rmilewi@sandia.gov

Darko Marniov

Univerity of Illinois Urbana Champaign
Urbana, IL
0000-0001-5023-3492

Abstract—Software tends to break or “collapse” over time, even if it is unchanged, due to non-obvious changes in the computational environment. Collapse in computational experiments undermines long-term credibility and hinders day-to-day operations. We propose to create the first public dataset which measures software collapse in computational experiments. We explain how that data can be used to identify best practices, make continuous testing feasible, and repair broken programs, which can greatly improve reproducibility of computational experiments.

Index Terms—reproducibility; software reliability; continuous testing

I. INTRODUCTION

Software tends to break over time, even if it is unchanged, due to non-obvious changes in the computational environment. This phenomenon is called “software collapse” [1], because software with an unstable foundation is analogous to a building with unstable foundation.¹ Software collapse is not a significant problem in some domains; it is acceptable if Google returns slightly different results one day to the next. But in the scientific domain, software collapse could manifest as irreproducible² results which hinder not only the long-term credibility of science but also its day-to-day operations.

1. **Hinders long-term credibility:** More than half of scientists surveyed across all fields develop software for their research [4]. If computational experiments are allowed to collapse, scientists cannot independently verify or build on each others’ results. This undermines two fundamental norms of science identified by Merton, organized skepticism and com-

munalism [5], that make science self-correcting. In recent years, this has manifested itself as the ongoing reproducibility crisis in computational science [6], which damages the long-term credibility of science [7].

2. **Hinders day-to-day operations:** Consider scientists tasked with securing their nations’ nuclear stockpile. They might create a simulation that tests if a physical part is going to properly preform a critical function for nuclear storage. The physical part might last several decades, but the software often collapses much faster than that. As our understanding of material science improves, they might want to reassess if the simulation still predicts the part preforms its function properly given our improved understanding. If the simulation experienced software collapse, this may be extremely difficult or impossible, especially if the original developer is retired.

Unfortunately, software collapse seems widespread. Zhao et al. studied software collapse computational of experiments deposited in the myExperiment registry [8]. They found that 80% of the experiments in their selection did not work, for a variety of causes: change of third-party resources, unavailable example data, insufficient execution environment, and insufficient metadata. Of these, change of third-party resources caused the most failures. This included a step in the experiment that referenced data from another server through the internet which was no longer available.

While the problem of irreproducible science is not solely technical, this paper studies technical solutions which should be a part of a holistic effort, including policy, economic and social factors. The technical solution could be proactive or reactive: a proactive solution would change something about the environment or application to provide determinism, whereas a reactive solution would seek to detect non-determinism and alert human developers. Proactive solutions are preferred because they offer certain

¹DSK: and all scientific software has an unstable foundation, unless it’s monolithic.

²In this article, we use Claerbout’s terminology [2]. “Reproducibility” means that one can use the same code in a different computational environment to get the same result [DSK: this definition is a bit loose]. If the execution does not terminate successfully, one can consider the error as “the result”; if code crashes in one environment and succeeds in another, that would count as an irreproducibility. Reproducibility is called “replicability” by some authors; see Plesser [3] for a discussion of terminology.

guarantees of determinism, although this guarantee is often not met in practice.

- **Snapshotting the environment:** Container images (e.g., Docker, qcow2), VM images, CDE [9], and Sumatra [10] attempt to snapshot the entire computational environment. Then, one can ship the entire filesystem to another user so they can reproduce the execution. However, this approach is heavyweight with filesystem snapshots as large as 50 Gb, as it needs to record a large chunk of the filesystem. Finally, these are difficult to modify and audit.
- **Specifying construction of the environment:** Dockerfiles let the user specify instructions to construct the computational environment enclosing software. However, these instructions are UNIX commands, which can be non-deterministic themselves³, e.g. `pip install`.
- **Functional package managers:** Functional package managers (e.g., Nix, Guix, Spack) is a restricted form of environment construction which only permits certain UNIX capabilities. For example, Nix only lets users download from the internet if they provide a hash of the expected outcome; if this hash is different, Nix errors because this execution is not going to reproduce the previous one. While this is useful for setting up the environment, this approach is only applied to the installation phase not the actual execution phase. It would be too burdensome to provide hash for every network access within the application, if the application talks to a database for example.

[DSK: maybe talk about version pinning?]

Furthermore, none of these proactive solutions can mitigate non-determinism due to network resources, pseudorandomness, and parallel program order. Zhao et al. showed that first of these, networked resources, is the most common cause of software collapse as well [8]. This is empirically validated, as Henkel et al. find 25% of Dockerfiles in their already limited sample still fail to build [11]. Therefore, we need both proactive *and* reactive solutions.

- **Continuous testing (reactive)**⁴ Continuous testing is a reactive solution to software collapse that is robust to networked resources, pseudorandomness, and parallel program order. ⁵ One could imagine running the computational experiment periodically

³Docker itself never claims that `Dockerfiles` are reproducible; it only says “reproducible” three times in [their documentation](#) at the time of writing, and none of them are referring to reproducing the same result from running a `Dockerfile` twice.

⁴3

⁵DSK: I’m not sure about the parallel program order part. SAG: this can be done by injecting randomness into the program schedule; recent work does something like that. I have yet to go back and find that paper, and weave it into the narrative in this paragraph.

to assess if the experiment is both not crashing and still producing the same results. The major drawback is increased computational cost. However, one can always lower the frequency of testing, which trades off computational resources with efficacy of finding bugs. Additionally, one could test mission-critical experiments more frequently than other experiments. If one could predict which workflows were more likely to break, one could also prioritize testing on that basis.

⁶: SAG: Explain difference with traditional CI.




Figure 1: Predicting the rate of software collapse can reduce the resource utilization and increase efficacy of continuous testing.

Hinsen suggests that most code should build on reliable, well-tested libraries can provide some degree of resistance to collapse [1]. In practice, many experiments fall into collapse despite their best effort to build on reliable foundations. If that level of reliability is insufficient, one can add continuous testing to help get more reliability.

This paper will build a dataset of software collapse of computational experiments and answer the following questions: 1. What are typical rates of decay? This number is not well-known, since the last experiment to measure it was Zhao et al., and we have new reproducibility technology (NextFlow over Taverna). 2. Can we predict the rate of decay for a project based on its history (if available) and code? A predictive model is important for the next research question. The model should function on a “cold start”, where we know nothing about the computational experiments historical results, but it should be able to learn from those historical runs if they are present. 3. Can we improve the efficiency of continuous testing by predicting the rate of decay? This could be useful for institutions, such as national labs, wanting to ensure their computational experiments remain valid while using resources efficiently. 4. What are the best practices that improve reproducibility? This lets

us make recommendations that are empirically backed. 5. In what fraction of the cases does automatic repair work? Automatic repair could let one run old workflows off-the-shelf with no modification.

II. METHODS (COLLECTING DATA)

We plan to collect data on software collapse of computational experiments by automatically running computational experiments from public registries. These registries include:

- **nf-core:** *TODO: describe each of these (one sentence).*
- [Dockstore](#)
- [Snakemake Catalog](#)
- [WorkflowHub](#)
- [myExperiment](#)
- [PegasusHub](#)
- Sandia’s internal repository
- Globus Flows
 - <https://www.globus.org/platform/services/flows>
 - <https://anl-braid.github.io/braid/>

We cannot take one computational experiment and simulate it one, five, and ten years into the future. Instead, we will look for historical revisions⁷ of an experiment from one, five, or ten years ago and simulate it today. All of the registries above store historical revisions⁸ of the workflow. We make a *time symmetry* assumption: historical rates of change will be similar to the future rate of change. It is likely that some will still work and some will fail, due to software collapse.

We will run the following pseudo-code to collect the data. Then we will analyze it as described in the next section. Finally, we plan to publish the raw data we collect for other researchers.

```
for registry in registries:
    for experiment in registry:
        for revision in experiment:
            for i in range(num_repetitions):
                execution = execute(revision)
                data.append((
                    execution.date,    execution.output,
                    execution.logs,    execution.resource_utilization,
                    revision.date,     revision.code,
                    experiment.name,   registry.name,
                ))
```

III. ANALYSIS

- **RQ1: measure rate of collapse:** We plan to replicate the quantities described by Zhao et al. [8] to see if these have changed since that work, or if they are different for workflows⁹: proportion of broken

experiments, and proportion of breakages due to each reason (volatile third-party resources, missing example data, missing execution environment, insufficient description). To this, we add “reproducible results” as a new “level” of success, beyond merely not crashing. We also plan to extend the failure classification of Zhao et al. by going into deeper subcategories. We will also study how the proportion of broken experiments changes with time.

- **RQ2: predict rate of collapse:** We will develop predictive models based on the history of failures, staleness, properties of the code in the revision, and other determinants to predict the probability that a given experiment will fail. We will use information theory criteria to quantify the difference from our predicted distribution to the actual distribution.
- **RQ3: improve continuous testing:** We can improve resource utilization of continuous testing by using our dataset to predict the rate of collapse of various computational experiments. Testing experiments prone to failure more often than reliable ones could save computational resources while maintaining approximately the same degree of reliability in all experiments.
- **RQ4: identify best practices:** We can also use this data to identify practices that improve the reproducibility and longevity of computational experiments. We will use a “Bayes net” to test for confounding causal variables.
- **RQ5: attempt automatic repair:** Once we know what kinds of failure are possible, we can also investigate automatic repair. Our dataset will contain the output logs for each failure. Therefore, we can apply similar techniques to Shipwright [11], such as using a language model to categorize many failures into a few clusters.

A. Threats to Validity

There are a number of threats to the validity of our work and planned results.

1. Our time symmetry assumption may not hold. With contemporary efforts on reproducibility, future rates of change may be markedly less than past rates of change. While our computed rates of change will be underestimates, those underestimates can still be useful as bounds. Our method will also be useful, unchanged, for future studies.
2. It is possible that our sample is not representative of the real world of computational experiments. However, we are casting the widest net we can by systematically pulling many experiments from several registries. Still, there is a selection bias in which workflows end up in registries. The model has some factors based on the population and some based on the actual history of the experiment. Its initial guess

⁷DSK: versions?

⁸DSK: versions?

⁹SAG: Let’s simplify this

when there is no history would be biased by our selection, but in the long run it would learn the characteristics of the actual experiment.

IV. CONCLUSION

Software collapse is an important yet understudied problem. We don't know the rate of software collapse in contemporary computational experiments. In order to do any research in this area, we need to build standard, communal datasets on software collapse.

The dataset would indicate how various reactive solutions compare, allowing us to identify the best practices that correlate with reproducibility. However, no proactive solution is perfect, so we also look at the reactive solution of continuous testing. The dataset would also allow us to optimize continuous testing such that it is feasible. Finally, when the continuous testing finds a failure, the automatic repair we plan to prototype would help fix that failure.

V. REFERENCES

- [1] K. Hinsén, "Dealing With Software Collapse," *Computing in Science & Engineering*, vol. 21, no. 3, pp. 104–108, May 2019, doi: [10.1109/MCSE.2019.2900945](https://doi.org/10.1109/MCSE.2019.2900945).
- [2] J. F. Claerbout and M. Karrenbach, "Electronic documents give reproducible research a new meaning," in *SEG Technical Program Expanded Abstracts 1992*, Jan. 1992, pp. 601–604, doi: [10.1190/1.1822162](https://doi.org/10.1190/1.1822162) [Online]. Available: <http://library.seg.org/doi/abs/10.1190/1.1822162>. [Accessed: Jun. 01, 2022]
- [3] H. E. Plesser, "Reproducibility vs. Replicability: A Brief History of a Confused Terminology," *Frontiers in Neuroinformatics*, vol. 11, 2018 [Online]. Available: <https://www.frontiersin.org/articles/10.3389/fninf.2017.00076>. [Accessed: Oct. 11, 2022]
- [4] S. Hettrick, "Software-saved/Software_in_research_survey_2014: Software In Research Survey." Zenodo, Feb. 2018 [Online]. Available: <https://zenodo.org/record/1183562>. [Accessed: May 26, 2022]
- [5] R. K. Merton, *The sociology of science: Theoretical and empirical investigations*, 4. Dr. Chicago: Univ. of Chicago Pr, 1974.
- [6] C. Collberg and T. A. Proebsting, "Repeatability in computer systems research," *Communications of the ACM*, vol. 59, no. 3, pp. 62–69, Feb. 2016, doi: [10.1145/2812803](https://doi.org/10.1145/2812803). [Online]. Available: <https://dl.acm.org/doi/10.1145/2812803>. [Accessed: May 27, 2022]
- [7] S. Ritchie, *Science Fictions: How Fraud, Bias, Negligence, and Hype Undermine the Search for Truth*, Illustrated edition. New York: Metropolitan Books, 2020.
- [8] J. Zhao *et al.*, "Why workflows break — understanding and combating decay in Taverna workflows," in *2012 IEEE 8th International Conference on E-Science (e-Science)*, Oct. 2012, p. 9, doi: [10.1109/eScience.2012.6404482](https://doi.org/10.1109/eScience.2012.6404482) [Online]. Available: [https://www.research.manchester.ac.uk/portal/en/publications/why-workflows-break--understanding-and-combating-decay-in-taverna-workflows\(cb-a81ca4-e92c-408e-8442-383d1f15fcdcf\)/export.html](https://www.research.manchester.ac.uk/portal/en/publications/why-workflows-break--understanding-and-combating-decay-in-taverna-workflows(cb-a81ca4-e92c-408e-8442-383d1f15fcdcf)/export.html)
- [9] P. Guo and D. Engler, "CDE: Using System Call Interposition to Automatically Create Portable Software Packages," in *2011 USENIX Annual Technical Conference*, Jun. 2011 [Online]. Available: https://www.usenix.org/legacy/events/atc11/tech/final_files/GuoEngler.pdf
- [10] A. Davison, "Automated Capture of Experiment Context for Easier Reproducibility in Computational Research," *Computing in Science & Engineering*, vol. 14, no. 4, pp. 48–56, Jul. 2012, doi: [10.1109/MCSE.2012.41](https://doi.org/10.1109/MCSE.2012.41). [Online]. Available: <http://ieeexplore.ieee.org/document/6180156/>. [Accessed: Jul. 08, 2022]
- [11] J. Henkel, D. Silva, L. Teixeira, M. d'Amorim, and T. Reps, "Shipwright: A Human-in-the-Loop System for Dockerfile Repair," in *2021 IEEE/ACM 43rd International Conference on Software Engineering (ICSE)*, May 2021, pp. 1148–1160, doi: [10.1109/ICSE43902.2021.00106](https://doi.org/10.1109/ICSE43902.2021.00106).